

# **BLG 317E**

# **DATABASE SYSTEMS**

## **WEEK 2**

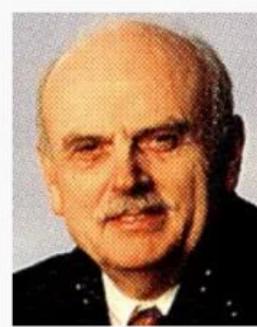
Slides credit:

S. Shivakumar, CS 145, Stanford University  
Database System Concepts, Silberschatz, Korth, Sudarshan

# Intro to Relational Model

# Relational Model

- Relational model is the most widely used database model.
- It is implemented by many commercial and open source vendors (companies).
- SQL (Structured Query Language) was developed by IBM in 1970s.



1969: Edgar F. "Ted" Codd invents the relational database.

**1969:** Edgar F. "Ted" Codd invents the relational database.

**1973:** Cullinane, led by John J. Cullinane, ships IDMS, a network-model database for IBM mainframes.

**1976:** Honeywell ships Multics Relational Data Store, the first commercial relational database.

**1979:** Oracle introduces the first commercial SQL relational database management system.

**1983:** IBM introduces DB2.

**1985:** The first business intelligence system is designed for Procter & Gamble.

**1991:** W.H. "Bill" Inmon publishes Building the Data Warehouse.



1991: W.H. "Bill" Inmon publishes Building the Data Warehouse.



# A Motivating, Running Example

Consider building a course management system (CMS):

*Entities* (e.g., Students, Courses)

*Relationships* (e.g., Ahmet is enrolled in CS 350)





## Students



File Edit View Insert Format Data Tools Add-ons Help All changes saved in Drive

undo redo print preview 100% \$ % .0 .00 123 Arial 10 B I S A

|    | A      | B                 | C                  | D              | E                 | F | G |
|----|--------|-------------------|--------------------|----------------|-------------------|---|---|
| 1  |        |                   |                    |                |                   |   |   |
| 2  |        |                   |                    |                |                   |   |   |
| 3  |        |                   |                    |                |                   |   |   |
| 4  |        |                   |                    |                |                   |   |   |
| 5  |        | <b>Student</b>    | <b>SID</b>         | <b>Address</b> |                   |   |   |
| 6  | Mickey |                   | 40001              | 43 Toontown    |                   |   |   |
| 7  | Daffy  |                   | 40002              | 147 Main St    |                   |   |   |
| 8  | Donald |                   | 50003              | 312 Escondido  |                   |   |   |
| 9  | Minnie |                   | 50004              | 451 Gates      |                   |   |   |
| 10 | Pluto  |                   | 10008              | 97 Packard     |                   |   |   |
| 11 |        |                   |                    |                |                   |   |   |
| 12 |        |                   |                    |                |                   |   |   |
| 13 |        |                   |                    |                |                   |   |   |
| 14 |        | <b>Course</b>     | <b>Description</b> | <b>Room</b>    | <b>Class size</b> |   |   |
| 15 | cs145  | Toon systems      | Nvidia             |                | 300               |   |   |
| 16 | cs161  | Animation art     | Gates 300          |                | 145               |   |   |
| 17 | cs245  | Painting town rec | Packard 45         |                | 27                |   |   |
| 18 |        |                   |                    |                |                   |   |   |

# 'Modeling' the CMS

## Logical Schema

Students(sid: *string*, name: *string*, gpa: *float*)

Courses(cid: *string*, cname: *string*, credits: *int*)

Enrolled(sid: *string*, cid: *string*, grade: *string*)

| sid | Name | Gpa |
|-----|------|-----|
| 101 | Bob  | 3.2 |
| 123 | Mary | 3.8 |

Students

Corresponding  
keys

| sid | cid | Grade |
|-----|-----|-------|
| 123 | 564 | A     |

Enrolled

| cid | cname | credits |
|-----|-------|---------|
| 564 | 564-2 | 4       |
| 308 | 417   | 2       |

Courses





# Set algebra (reminder)

List: [1, 1, 2, 3]

Set: {1, 2, 3}

Multiset: {1, 1, 2, 3}

A **multiset** is an unordered list (or: a set with multiple duplicate instances allowed)

## UNIONS

Set: {1, 2, 3}  $\cup$  { 2 } = { 1, 2, 3 }

Multiset: {1, 1, 2, 3}  $\cup$  { 2 } = { 1, 1, 2, 2, 3 }



# Tables

## Product

| PName       | Price    | Manuf      |
|-------------|----------|------------|
| Gizmo       | \$19.99  | GizmoWorks |
| Powergizmo  | \$29.99  | GizmoWorks |
| SingleTouch | \$149.99 | Canon      |
| MultiTouch  | \$203.99 | Hitachi    |

A relation or table is a multiset of tuples having the attributes specified by the schema

*Let's break this definition down*



# Tables

Product

| PName       | Price    | Manuf      |
|-------------|----------|------------|
| Gizmo       | \$19.99  | GizmoWorks |
| Powergizmo  | \$29.99  | GizmoWorks |
| SingleTouch | \$149.99 | Canon      |
| MultiTouch  | \$203.99 | Hitachi    |

An attribute (or column or field) is a typed data entry present in each tuple in the relation

*Attributes must have an atomic type in standard SQL, i.e. not a list, set, etc.*

STUDENT Table

| student_id | name  | gpa | courses_taken          |
|------------|-------|-----|------------------------|
| 53410      | Jones | 3.4 | Mat101, Kim101, Fiz101 |

Standard SQL does not have array datatype.  
(PostgreSQL DBMS supports arrays.)



# Tables

## Product

| PName       | Price    | Manuf      |
|-------------|----------|------------|
| Gizmo       | \$19.99  | GizmoWorks |
| Powergizmo  | \$29.99  | GizmoWorks |
| SingleTouch | \$149.99 | Canon      |
| MultiTouch  | \$203.99 | Hitachi    |

A **tuple** or **row** is a single entry in the table having the attributes specified by the schema

Also referred to sometimes as a **Record**



# Tables

## Product

| PName       | Price    | Manuf      |
|-------------|----------|------------|
| Gizmo       | \$19.99  | GizmoWorks |
| Powergizmo  | \$29.99  | GizmoWorks |
| SingleTouch | \$149.99 | Canon      |
| MultiTouch  | \$203.99 | Hitachi    |

The number of tuples is the cardinality of the relation

The number of attributes is the arity/degree of the relation

Tuple order is insignificant.  
Attribute order is insignificant.



# Data Types (Domains)

Atomic types:

Characters: CHAR(20), VARCHAR(50)

Numbers: INT, BIGINT, SMALLINT, FLOAT

Others: MONEY, DATETIME...

Every attribute must have an atomic type

Hence tables are flat



# Table Schemas

The **schema** of a table is the table name, its attributes, and their types:

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

A **key** is an attribute or a set of attributes whose values are unique; we underline a key.

Product(Pname: *string*, Price: *float*, Category: *string*, Manufacturer: *string*)

# Example of a Relation

The diagram illustrates a relation table with four columns: ID, name, dept\_name, and salary. The first column, ID, contains numerical values. The second column, name, contains names. The third column, dept\_name, contains department names. The fourth column, salary, contains monetary values. Arrows point from the text labels 'attributes (or columns)' to the column headers, and arrows point from the text labels 'tuples (or rows)' to the data rows.

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 12121     | Wu          | Finance          | 90000         |
| 15151     | Mozart      | Music            | 40000         |
| 22222     | Einstein    | Physics          | 95000         |
| 32343     | El Said     | History          | 60000         |
| 33456     | Gold        | Physics          | 87000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 58583     | Califieri   | History          | 62000         |
| 76543     | Singh       | Finance          | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |

# Relational Data Model - Summary

- Relational database contains a set of relations (tables).
- A relation is a table with rows and columns.
- Every relation has a schema definition, which describes the columns (fields).
- **Schema** : Specifies name of relation (table), also names and data types of each column in table.
- **Instance** : A table with rows (records, tuples) and columns (fields, attributes), containing the actual data.
  - Number of rows : cardinality
  - Number of fields : arity/degree

# Example: **PUBLICATIONS Database Schema**

- The following is a conceptual schema definition.  
(They are not SQL commands, only informal definitions.)
- Table names and field names are written.
- Data types of fields may be omitted.
- Primary key fields are underlined.

## Database Schema

```
PUBLISHER (publisher_id, publisher_name, phone)  
AUTHOR      (author_id, author_name, email)  
BOOK        (ISBN, publisher_id, author_id, title)
```

# Key Attributes

## Key Attribute:

- must **uniquely** identify a row in a table.
- No duplicates, no null values.
- Example: **author\_id** column

## Candidate Key:

- Each alternative key is called a candidate key.
- One of the candidate keys is arbitrarily designated to be the primary key.
- Example: **email column** is a candidate key.

**Attributes  
(columns / fields)**

| author_id | author_name  | email    |
|-----------|--------------|----------|
| 100       | Elmasri      | elm@uta  |
| 200       | Ullman       | ulm@ibm  |
| 300       | Weinberg     | null     |
| 400       | Silberschatz | slb@ms   |
| 500       | Ramakrishnan | rma@wisc |

**Tuples  
(rows / records)**

# Primary Key

- A Primary Key is a unique attribute, or a group of multiple attributes, that uniquely identify a record in a table without duplications.
- If there are more than one **candidate** key in a table, one is selected as the primary key, others are alternate keys.
- Any attribute that is part of the primary key can not be empty in any tuple.
- Every relation (table) should have a primary key.

# Surrogate (Substitute) Key

- If a natural key can not be found, a surrogate (substitute) key can be defined.
- It can be generated by the DBMS system as an **auto-incremented** counter.

AUTHOR Table

| <u>author_id</u> | author_name  | email    |
|------------------|--------------|----------|
| 1                | Elmasri      | elm@uta  |
| 2                | Ullman       | ulm@ibm  |
| 3                | Weinberg     | null     |
| 4                | Silberschatz | slb@ms   |
| 5                | Ramakrishnan | rma@wisc |



Surrogate Key:

Sequential record numbers are (record\_number) automatically assigned by DBMS, every time a new record is added to table.

# Super Key

- Let  $K \subseteq R$
- $K$  is a **super key** of  $R$  if values for  $K$  are sufficient to identify a unique tuple of each possible relation  $r(R)$ 
  - Example:  $\{ID\}$  and  $\{ID, name\}$  are both superkeys of *instructor*.
- Superkey  $K$  is a **candidate key** if  $K$  is minimal  
Example:  $\{ID\}$  is a candidate key for *Instructor*

# Foreign Key

- **Foreign Key:** An attribute of a table that is a Candidate key of another table.
- Foreign keys are for implementation of relations (associations) between tables.
- Must correspond (reference) to the primary/candidate key of the other table.
- DBMS enforces all foreign key constraints, to maintain **referential integrity**.

**Primary Key**

## PUBLISHER table

| <u>publisher_id</u> | <u>publisher_name</u> | phone         |
|---------------------|-----------------------|---------------|
| 10                  | Addison-Wesley        | 212-303-8421  |
| 20                  | Pearson Prentice Hall | null          |
| 30                  | McGraw-Hil            | 216-404-02354 |

## BOOK table

**Foreign Key      Foreign Key**

| <u>ISBN</u> | <u>publisher_id</u> | <u>author_id</u> | title                     |
|-------------|---------------------|------------------|---------------------------|
| A0012345    | 10                  | 100              | Fundamentals of Databases |
| B0033224    | 20                  | 200              | Database Systems          |
| C0092831    | 30                  | 300              | SQL Reference             |
| D0027354    | 30                  | 400              | Database Concepts         |
| E0091237    | 30                  | 500              | Database Management       |
| F0072261    | 20                  | 100              | Database Design           |

# Foreign Keys

- *dept\_name* in *instructor* is a foreign key from *instructor* referencing *department*

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222     | Einstein    | Physics          | 95000         |
| 12121     | Wu          | Finance          | 90000         |
| 32343     | El Said     | History          | 60000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 58583     | Califieri   | History          | 62000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 15151     | Mozart      | Music            | 40000         |
| 33456     | Gold        | Physics          | 87000         |
| 76543     | Singh       | Finance          | 80000         |

(a) The *instructor* table

| <i>dept_name</i> | <i>building</i> | <i>budget</i> |
|------------------|-----------------|---------------|
| Comp. Sci.       | Taylor          | 100000        |
| Biology          | Watson          | 90000         |
| Elec. Eng.       | Taylor          | 85000         |
| Music            | Packard         | 80000         |
| Finance          | Painter         | 120000        |
| History          | Painter         | 50000         |
| Physics          | Watson          | 70000         |

(b) The *department* table

# Referential Integrity

- **Referential Integrity:** All values of a **foreign key** attribute in a table must be present, among the values of the referenced table's **primary key** or candidate key attribute.
- If there is no corresponding Primary Key for a Foreign Key, a runtime error occurs.
- If a request (Add/Delete/Update) would break referential integrity, the DBMS can use one of the following actions:
  - Don't allow, and return an error (default action of DBMS)
  - Reflect the change to affected tuples
  - Assign null value to foreign key
  - Assign default value to foreign key

# Referential Integrity

## Example:

- Insert operation for the following row is **not allowed** by DBMS.
  - Because **publisher\_id** and **author\_id** must exist in the related tables (PUBLISHER and AUTHOR).

**BOOK Table**

| <u>ISBN</u> | <u>publisher_id</u> | <u>author_id</u> | title           |
|-------------|---------------------|------------------|-----------------|
| A0036721    | 0                   | 0                | Digital Systems |

# Relational Query Languages

- Procedural versus non-procedural, or declarative
- “Pure” languages:
  - Relational algebra
  - Tuple relational calculus
  - Domain relational calculus
- The above 3 pure languages are equivalent in computing power
- We will concentrate in this semester only on relational algebra
  - Procedural language consisting of a set of operations that take one or two relations as input and produce a new relation as their result.
  - Not turing-machine equivalent
    - ▶ (not equal to a programming language in expressive power)

# Relational Algebra

## Basic Operations

| Operation     | Symbol    | Description                                                                                                   |
|---------------|-----------|---------------------------------------------------------------------------------------------------------------|
| Selection     | $\sigma$  | Selects a subset of rows from relation (horizontal)                                                           |
| Projection    | $\pi$     | Retains only wanted columns from relation (vertical).                                                         |
| Cross-product | $\times$  | Allows to combine each row of the left relation with the each row of the right relations (Cartesian product). |
| Join          | $\bowtie$ | A compound operator involving cross product, selection, and projection.                                       |

# Relational Algebra

## Set Operations

(R1 and R2 are tables.)

| Operation    | Symbol | Description                  |
|--------------|--------|------------------------------|
| Union        | $\cup$ | Tuples in R1 or in R2.       |
| Intersection | $\cap$ | Tuples in R1 and in R2.      |
| Difference   | $-$    | Tuples in R1, but not in R2. |

# Select Operation

- The **select** operation selects tuples that satisfy a given predicate.
- Notation:  $\sigma_p(r)$
- $p$  is called the **selection predicate**
- Example: select those tuples of the *instructor* relation where the instructor is in the “Physics” department.
  - Query

$$\sigma_{dept\_name = "Physics"}(instructor)$$

- Result

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 22222 | Einstein   | Physics    | 95000  |
| 12121 | Wu         | Finance    | 90000  |
| 32343 | El Said    | History    | 60000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 58583 | Califieri  | History    | 62000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 76543 | Singh      | Finance    | 80000  |



| ID    | name     | dept_name | salary |
|-------|----------|-----------|--------|
| 22222 | Einstein | Physics   | 95000  |
| 33456 | Gold     | Physics   | 87000  |

# Select Operation (Cont.)

- We allow comparisons using  
 $=, \neq, >, \geq, <, \leq$   
in the selection predicate.
- We can combine several predicates into a larger predicate by using the connectives:  
 $\wedge$  (**and**),  $\vee$  (**or**),  $\neg$  (**not**)
- Example: Find the instructors in Physics with a salary greater \$90,000, we write:

$$\sigma_{dept\_name = "Physics"} \wedge salary > 90,000 (instructor)$$

- Besides, select predicate may include comparisons between two attributes.
  - Example, find all departments whose name is the same as their building name:
    - $\sigma_{dept\_name=building} (department)$

# Project Operation

- A unary operation that returns its argument relation, with certain attributes left out.
- Notation:

$$\prod_{A_1, A_2, A_3 \dots A_k} (r)$$

where  $A_1, A_2$  are attribute names and  $r$  is a relation name.

- The result is defined as the relation of  $k$  columns obtained by erasing the columns that are not listed

# Project Operation (Cont.)

- Example: eliminate the *dept\_name* attribute of *instructor*
- Query:

$$\Pi_{ID, name, salary} (instructor)$$

- Result:

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222     | Einstein    | Physics          | 95000         |
| 12121     | Wu          | Finance          | 90000         |
| 32343     | El Said     | History          | 60000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |
| 76766     | Crick       | Biology          | 72000         |
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 58583     | Califieri   | History          | 62000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 15151     | Mozart      | Music            | 40000         |
| 33456     | Gold        | Physics          | 87000         |
| 76543     | Singh       | Finance          | 80000         |



| <i>ID</i> | <i>name</i> | <i>salary</i> |
|-----------|-------------|---------------|
| 10101     | Srinivasan  | 65000         |
| 12121     | Wu          | 90000         |
| 15151     | Mozart      | 40000         |
| 22222     | Einstein    | 95000         |
| 32343     | El Said     | 60000         |
| 33456     | Gold        | 87000         |
| 45565     | Katz        | 75000         |
| 58583     | Califieri   | 62000         |
| 76543     | Singh       | 80000         |
| 76766     | Crick       | 72000         |
| 83821     | Brandt      | 92000         |
| 98345     | Kim         | 80000         |

# Composition of Relational Operations

- The result of a relational-algebra operation is relation.
- Therefore, relational-algebra operations can be composed together into a **relational-algebra expression**.

**QUERY : Get author names and their email address, whose ID number is less than 300.**

**RA Query**

$$\pi_{\text{author\_name}, \text{email}} (\sigma_{\text{author\_id} < 300} (\text{AUTHOR}))$$

Projection operation  
is applied second.

Selection operation  
is applied first.

| author_id | author_name  | email    |
|-----------|--------------|----------|
| 100       | Elmasri      | elm@uta  |
| 200       | Ullman       | ulm@ibm  |
| 300       | Weinberg     | null     |
| 400       | Silberschatz | slb@ms   |
| 500       | Ramakrishnan | rma@wisc |



| author_name | email   |
|-------------|---------|
| Elmasri     | elm@uta |
| Ullman      | ulm@ibm |

# Cartesian-Product Operation

- The Cartesian-product operation (denoted by  $\times$ ) allows us to combine information from any two relations.
  - Example: the Cartesian product of the relations *instructor* and *teaches* is written as:

*instructor*  $\times$  *teaches*
- We construct a tuple of the result out of each possible pair of tuples:
  - one from the *instructor* relation and one from the *teaches* relation (see the next slides)
- Since the instructor *ID* appears in both relations we distinguish between these attributes by attaching to the attribute the name of the relation from which the attribute originally came.
  - *instructor.ID*
  - *teaches.ID*

|           |                  |                  |                 |
|-----------|------------------|------------------|-----------------|
| <i>ID</i> | <i>name</i>      | <i>dept_name</i> | <i>salary</i>   |
| <i>ID</i> | <i>course_id</i> | <i>sec_id</i>    | <i>semester</i> |

# The *instructor* x *teaches* table

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 22222 | Einstein   | Physics    | 95000  |
| 12121 | Wu         | Finance    | 90000  |
| 32343 | El Said    | History    | 60000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 58583 | Califieri  | History    | 62000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 76543 | Singh      | Finance    | 80000  |

X

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |
| 32343 | HIS-351   | 1      | Spring   | 2010 |
| 45565 | CS-101    | 1      | Spring   | 2010 |
| 45565 | CS-319    | 1      | Spring   | 2010 |
| 76766 | BIO-101   | 1      | Summer   | 2009 |
| 76766 | BIO-301   | 1      | Summer   | 2010 |
| 83821 | CS-190    | 1      | Spring   | 2009 |
| 83821 | CS-190    | 2      | Spring   | 2009 |
| 83821 | CS-319    | 2      | Spring   | 2010 |
| 98345 | EE-181    | 1      | Spring   | 2009 |



| instructor.ID | name       | dept_name  | salary | teaches.ID | course_id | sec_id | semester | year |
|---------------|------------|------------|--------|------------|-----------|--------|----------|------|
| 10101         | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-101    | 1      | Fall     | 2017 |
| 10101         | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-315    | 1      | Spring   | 2018 |
| 10101         | Srinivasan | Comp. Sci. | 65000  | 10101      | CS-347    | 1      | Fall     | 2017 |
| 10101         | Srinivasan | Comp. Sci. | 65000  | 12121      | FIN-201   | 1      | Spring   | 2018 |
| 10101         | Srinivasan | Comp. Sci. | 65000  | 15151      | MU-199    | 1      | Spring   | 2018 |
| 10101         | Srinivasan | Comp. Sci. | 65000  | 22222      | PHY-101   | 1      | Fall     | 2017 |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 12121         | Wu         | Finance    | 90000  | 10101      | CS-101    | 1      | Fall     | 2017 |
| 12121         | Wu         | Finance    | 90000  | 10101      | CS-315    | 1      | Spring   | 2018 |
| 12121         | Wu         | Finance    | 90000  | 10101      | CS-347    | 1      | Fall     | 2017 |
| 12121         | Wu         | Finance    | 90000  | 12121      | FIN-201   | 1      | Spring   | 2018 |
| 12121         | Wu         | Finance    | 90000  | 15151      | MU-199    | 1      | Spring   | 2018 |
| 12121         | Wu         | Finance    | 90000  | 22222      | PHY-101   | 1      | Fall     | 2017 |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 15151         | Mozart     | Music      | 40000  | 10101      | CS-101    | 1      | Fall     | 2017 |
| 15151         | Mozart     | Music      | 40000  | 10101      | CS-315    | 1      | Spring   | 2018 |
| 15151         | Mozart     | Music      | 40000  | 10101      | CS-347    | 1      | Fall     | 2017 |
| 15151         | Mozart     | Music      | 40000  | 12121      | FIN-201   | 1      | Spring   | 2018 |
| 15151         | Mozart     | Music      | 40000  | 15151      | MU-199    | 1      | Spring   | 2018 |
| 15151         | Mozart     | Music      | 40000  | 22222      | PHY-101   | 1      | Fall     | 2017 |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| 22222         | Einstein   | Physics    | 95000  | 10101      | CS-101    | 1      | Fall     | 2017 |
| 22222         | Einstein   | Physics    | 95000  | 10101      | CS-315    | 1      | Spring   | 2018 |
| 22222         | Einstein   | Physics    | 95000  | 10101      | CS-347    | 1      | Fall     | 2017 |
| 22222         | Einstein   | Physics    | 95000  | 12121      | FIN-201   | 1      | Spring   | 2018 |
| 22222         | Einstein   | Physics    | 95000  | 15151      | MU-199    | 1      | Spring   | 2018 |
| 22222         | Einstein   | Physics    | 95000  | 22222      | PHY-101   | 1      | Fall     | 2017 |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |
| ...           | ...        | ...        | ...    | ...        | ...       | ...    | ...      | ...  |

# Cartesian-Product Operation

**Student**

| S_id | Name | Class | Age |
|------|------|-------|-----|
| 1    | Ali  | 5     | 25  |
| 2    | Veli | 10    | 30  |
| 3    | Ayse | 8     | 35  |

**Course**

| C_id | C_name        |
|------|---------------|
| 11   | Database Sys  |
| 21   | Operating Sys |

**Student X Course**

| S_id | Name | Class | Age | C_id | C_name        |
|------|------|-------|-----|------|---------------|
| 1    | Ali  | 5     | 25  | 11   | Database Sys  |
| 1    | Ali  | 5     | 25  | 21   | Operating Sys |
| 2    | Veli | 10    | 30  | 11   | Database Sys  |
| 2    | Veli | 10    | 30  | 21   | Operating Sys |
| 3    | Ayse | 8     | 35  | 11   | Database Sys  |
| 3    | Ayse | 8     | 35  | 21   | Operating Sys |

# Join Operation

- The Cartesian-Product

*instructor X teaches*

associates every tuple of instructor with every tuple of teaches.

- Most of the resulting rows have information about instructors who did NOT teach a particular course.
- To get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught, we write:

$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$

- We get only those tuples of “*instructor X teaches*” that pertain to instructors and the courses that they taught.
- The result of this expression, shown in the next slide

# Join Operation (Cont.)

- The table corresponding to:

$$\sigma_{instructor.id = teaches.id} (instructor \times teaches)$$

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 22222 | Einstein   | Physics    | 95000  |
| 12121 | Wu         | Finance    | 90000  |
| 32343 | El Said    | History    | 60000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 58583 | Califieri  | History    | 62000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 76543 | Singh      | Finance    | 80000  |



| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |
| 32343 | HIS-351   | 1      | Spring   | 2010 |
| 45565 | CS-101    | 1      | Spring   | 2010 |
| 45565 | CS-319    | 1      | Spring   | 2010 |
| 76766 | BIO-101   | 1      | Summer   | 2009 |
| 76766 | BIO-301   | 1      | Summer   | 2010 |
| 83821 | CS-190    | 1      | Spring   | 2009 |
| 83821 | CS-190    | 2      | Spring   | 2009 |
| 83821 | CS-319    | 2      | Spring   | 2010 |
| 98345 | EE-181    | 1      | Spring   | 2009 |

| ID    | name       | dept_name  | salary | course_id | sec_id | semester | year |
|-------|------------|------------|--------|-----------|--------|----------|------|
| 10101 | Srinivasan | Comp. Sci. | 65000  | CS-101    | 1      | Fall     | 2009 |
| 10101 | Srinivasan | Comp. Sci. | 65000  | CS-315    | 1      | Spring   | 2010 |
| 10101 | Srinivasan | Comp. Sci. | 65000  | CS-347    | 1      | Fall     | 2009 |
| 12121 | Wu         | Finance    | 90000  | FIN-201   | 1      | Spring   | 2010 |
| 15151 | Mozart     | Music      | 40000  | MU-199    | 1      | Spring   | 2010 |
| 22222 | Einstein   | Physics    | 95000  | PHY-101   | 1      | Fall     | 2009 |
| 32343 | El Said    | History    | 60000  | HIS-351   | 1      | Spring   | 2010 |
| 45565 | Katz       | Comp. Sci. | 75000  | CS-101    | 1      | Spring   | 2010 |
| 45565 | Katz       | Comp. Sci. | 75000  | CS-319    | 1      | Spring   | 2010 |
| 76766 | Crick      | Biology    | 72000  | BIO-101   | 1      | Summer   | 2009 |
| 76766 | Crick      | Biology    | 72000  | BIO-301   | 1      | Summer   | 2010 |
| 83821 | Brandt     | Comp. Sci. | 92000  | CS-190    | 1      | Spring   | 2009 |
| 83821 | Brandt     | Comp. Sci. | 92000  | CS-190    | 2      | Spring   | 2009 |
| 83821 | Brandt     | Comp. Sci. | 92000  | CS-319    | 2      | Spring   | 2010 |
| 98345 | Kim        | Elec. Eng. | 80000  | EE-181    | 1      | Spring   | 2009 |

# Join Operation (Cont.)

- The **join** operation allows us to combine a select operation and a Cartesian-Product operation into a single operation.
- Consider relations  $r (R)$  and  $s (S)$
- Let “theta” be a predicate on attributes in the schema R “union” S. The join operation  $r \bowtie_{\theta} s$  is defined as follows:

$$r \bowtie_{\theta} s = \sigma_{\theta}(r \times s)$$

- Thus

$$\sigma_{instructor.id = teaches.id}(instructor \times teaches))$$

- Can equivalently be written as

$$instructor \bowtie_{Instructor.id = teaches.id} teaches.$$

# Join Example

**Student**  $\bowtie_{\text{Student.Std} = \text{Course.Class}}$  **(Course)**

**Student**

| S_id | Name | Std | Age |
|------|------|-----|-----|
| 1    | Ali  | 5   | 25  |
| 2    | Veli | 10  | 30  |
| 3    | Ayse | 8   | 35  |

**Course**

| Class | C_name        |
|-------|---------------|
| 10    | Database Sys  |
| 5     | Operating Sys |

**Student**  $\bowtie_{\text{Student.Std} = \text{Course.Class}}$  **(Course)**

| S_id | Name | Std | Age | Class | C_name        |
|------|------|-----|-----|-------|---------------|
| 1    | Ali  | 5   | 25  | 5     | Operating Sys |
| 2    | Veli | 10  | 30  | 10    | Database Sys  |

# Join Types

- **Natural Join** : Uses attributes with the **same name**, that exists in both tables. Ex: *author*  $\bowtie$  *book*
  - The followings are intermediate steps.
    - Compute cross product R1 X R2
    - Select rows where attributes that appear in both relations have equal key values
    - Project all unique attributes and one copy of each of the common ones.
- **Equi-Join** : Special case of Natural join, where an equality condition operator is written on common columns (names may be different). Ex: *instructor*  $\bowtie$  *Instructor.id = teaches.id* *teaches*
- **Theta-Join** : Any join condition may be used. Ex:  $r \bowtie_{\theta} s$

# Join operation (Natural-Join)

QUERY : Join all author records with corresponding book records, over the common key field (**author ID**).

AUTHOR

| <u>author_id</u> | author_name | email |
|------------------|-------------|-------|
|------------------|-------------|-------|

BOOK

| <u>ISBN</u> | publisher_id | <u>author_id</u> | title |
|-------------|--------------|------------------|-------|
|-------------|--------------|------------------|-------|

RA Query

AUTHOR  $\bowtie$  BOOK

Result

| author_id | author_name  | email    | ISBN     | publisher_id | author_id | title                     |
|-----------|--------------|----------|----------|--------------|-----------|---------------------------|
| 100       | Elmasri      | elm@uta  | A0012345 | 10           | 100       | Fundamentals of Databases |
| 100       | Elmasri      | elm@uta  | F0072261 | 20           | 100       | Database Design           |
| 200       | Ullman       | ulm@ibm  | B0033224 | 20           | 200       | Database Systems          |
| 300       | Weinberg     | null     | C0092831 | 30           | 300       | SQL Reference             |
| 400       | Silberschatz | slb@ms   | D0027354 | 30           | 400       | Database Concepts         |
| 500       | Ramakrishnan | rma@wisc | E0091237 | 30           | 500       | Database Management       |

# Join operation (Natural-Join)

**QUERY : Join all author records with corresponding book records, over the common key field (author ID).**

AUTHOR

| <u>author_id</u> | author_name | email |
|------------------|-------------|-------|
|------------------|-------------|-------|

BOOK

| <u>ISBN</u> | publisher_id | <u>author_id</u> | title |
|-------------|--------------|------------------|-------|
|-------------|--------------|------------------|-------|

RA Query

AUTHOR  $\bowtie$  BOOK

Result

| author_id | author_name  | email    | <u>ISBN</u> | publisher_id | title                     |
|-----------|--------------|----------|-------------|--------------|---------------------------|
| 100       | Elmasri      | elm@uta  | A0012345    | 10           | Fundamentals of Databases |
| 100       | Elmasri      | elm@uta  | F0072261    | 20           | Database Design           |
| 200       | Ullman       | ulm@ibm  | B0033224    | 20           | Database Systems          |
| 300       | Weinberg     | null     | C0092831    | 30           | SQL Reference             |
| 400       | Silberschatz | slb@ms   | D0027354    | 30           | Database Concepts         |
| 500       | Ramakrishnan | rma@wisc | E0091237    | 30           | Database Management       |

Duplicate  
common  
column  
removed

# Union Operation

- The union operation allows us to combine two relations
  - Notation:  $r \cup s$
- For  $r \cup s$  to be valid.
  1.  $r, s$  must have the same **arity** (same number of attributes)
  2. The attribute domains of  $r$  and  $s$  must be **compatible**  
(example: 2<sup>nd</sup> column of  $r$  deals with the same type of values as does the 2<sup>nd</sup> column of  $s$ )
- Example: to find all courses taught in the Fall 2009 semester, or in the Spring 2010 semester, or in both

$$\Pi_{course\_id}(\sigma_{semester="Fall"} \wedge year=2009(section)) \cup$$
$$\Pi_{course\_id}(\sigma_{semester="Spring"} \wedge year=2010(section))$$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
|-----------|--------|----------|------|----------|-------------|--------------|

# Union Operation (Cont.)

- Result of:

$$\begin{aligned}\Pi_{course\_id}(\sigma_{semester="Fall"} \wedge year=2009(section)) \cup \\ \Pi_{course\_id}(\sigma_{semester="Spring"} \wedge year=2010(section))\end{aligned}$$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
| BIO-101   | 1      | Summer   | 2009 | Painter  | 514         | B            |
| BIO-301   | 1      | Summer   | 2010 | Painter  | 514         | A            |
| CS-101    | 1      | Fall     | 2009 | Packard  | 101         | H            |
| CS-101    | 1      | Spring   | 2010 | Packard  | 101         | F            |
| CS-190    | 1      | Spring   | 2009 | Taylor   | 3128        | E            |
| CS-190    | 2      | Spring   | 2009 | Taylor   | 3128        | A            |
| CS-315    | 1      | Spring   | 2010 | Watson   | 120         | D            |
| CS-319    | 1      | Spring   | 2010 | Watson   | 100         | B            |
| CS-319    | 2      | Spring   | 2010 | Taylor   | 3128        | C            |
| CS-347    | 1      | Fall     | 2009 | Taylor   | 3128        | A            |
| EE-181    | 1      | Spring   | 2009 | Taylor   | 3128        | C            |
| FIN-201   | 1      | Spring   | 2010 | Packard  | 101         | B            |
| HIS-351   | 1      | Spring   | 2010 | Painter  | 514         | C            |
| MU-199    | 1      | Spring   | 2010 | Packard  | 101         | D            |
| PHY-101   | 1      | Fall     | 2009 | Watson   | 100         | A            |



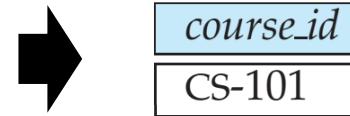
| course_id |
|-----------|
| CS-101    |
| CS-315    |
| CS-319    |
| CS-347    |
| FIN-201   |
| HIS-351   |
| MU-199    |
| PHY-101   |

# Set-Intersection Operation

- The set-intersection operation allows us to find tuples that are in both the input relations.
  - Notation:  $r \cap s$
- Assume:
  - $r, s$  have the *same arity*
  - attributes of  $r$  and  $s$  are compatible
- Example: Find the set of all courses taught in both the Fall 2009 and the Spring 2010 semesters.

$$\begin{array}{l} \prod_{course\_id} (\sigma_{semester="Fall"} \wedge year=2009 (section)) \cap \\ \prod_{course\_id} (\sigma_{semester="Spring"} \wedge year=2010 (section)) \end{array}$$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
| BIO-101   | 1      | Summer   | 2009 | Painter  | 514         | B            |
| BIO-301   | 1      | Summer   | 2010 | Painter  | 514         | A            |
| CS-101    | 1      | Fall     | 2009 | Packard  | 101         | H            |
| CS-101    | 1      | Spring   | 2010 | Packard  | 101         | F            |
| CS-190    | 1      | Spring   | 2009 | Taylor   | 3128        | E            |
| CS-190    | 2      | Spring   | 2009 | Taylor   | 3128        | A            |
| CS-315    | 1      | Spring   | 2010 | Watson   | 120         | D            |
| CS-319    | 1      | Spring   | 2010 | Watson   | 100         | B            |
| CS-319    | 2      | Spring   | 2010 | Taylor   | 3128        | C            |
| CS-347    | 1      | Fall     | 2009 | Taylor   | 3128        | A            |
| EE-181    | 1      | Spring   | 2009 | Taylor   | 3128        | C            |
| FIN-201   | 1      | Spring   | 2010 | Packard  | 101         | B            |
| HIS-351   | 1      | Spring   | 2010 | Painter  | 514         | C            |
| MU-199    | 1      | Spring   | 2010 | Packard  | 101         | D            |
| PHY-101   | 1      | Fall     | 2009 | Watson   | 100         | A            |



# Set Difference Operation

- The set-difference operation allows us to find tuples that are in one relation but are not in another.
- Notation  $r - s$
- Set differences must be taken between **compatible** relations.
  - $r$  and  $s$  must have the **same arity**
  - attribute domains of  $r$  and  $s$  must be compatible
- Example: to find all courses taught in the Fall 2009 semester, but not in the Spring 2010 semester

$$\prod_{course\_id} (\sigma_{semester='Fall' \wedge year=2009} (section)) - \\ \prod_{course\_id} (\sigma_{semester='Spring' \wedge year=2010} (section))$$

| course_id | sec_id | semester | year | building | room_number | time_slot_id |
|-----------|--------|----------|------|----------|-------------|--------------|
| BIO-101   | 1      | Summer   | 2009 | Painter  | 514         | B            |
| BIO-301   | 1      | Summer   | 2010 | Painter  | 514         | A            |
| CS-101    | 1      | Fall     | 2009 | Packard  | 101         | H            |
| CS-101    | 1      | Spring   | 2010 | Packard  | 101         | F            |
| CS-190    | 1      | Spring   | 2009 | Taylor   | 3128        | E            |
| CS-190    | 2      | Spring   | 2009 | Taylor   | 3128        | A            |
| CS-315    | 1      | Spring   | 2010 | Watson   | 120         | D            |
| CS-319    | 1      | Spring   | 2010 | Watson   | 100         | B            |
| CS-319    | 2      | Spring   | 2010 | Taylor   | 3128        | C            |
| CS-347    | 1      | Fall     | 2009 | Taylor   | 3128        | A            |
| EE-181    | 1      | Spring   | 2009 | Taylor   | 3128        | C            |
| FIN-201   | 1      | Spring   | 2010 | Packard  | 101         | B            |
| HIS-351   | 1      | Spring   | 2010 | Painter  | 514         | C            |
| MU-199    | 1      | Spring   | 2010 | Packard  | 101         | D            |
| PHY-101   | 1      | Fall     | 2009 | Watson   | 100         | A            |



# The Assignment Operation

- It is convenient at times to write a relational-algebra expression by assigning parts of it to temporary relation variables.
- The assignment operation is denoted by  $\leftarrow$  and works like assignment in a programming language.
- Example: Find all instructor in the “Physics” and Music department.

$$\text{Physics} \leftarrow \sigma_{\text{dept\_name} = \text{“Physics”}}(\text{instructor})$$
$$\text{Music} \leftarrow \sigma_{\text{dept\_name} = \text{“Music”}}(\text{instructor})$$
$$\text{Physics} \cup \text{Music}$$

- With the assignment operation, a query can be written as a sequential program consisting of a series of assignments followed by an expression whose value is displayed as the result of the query.

# Equivalent Queries

- There is more than one way to write a query in relational algebra.
- Example: Find instructors in the Physics department with salary greater than 90,000
  - Query 1

$$\sigma_{dept\_name = "Physics"} \wedge salary > 90,000 (instructor)$$

- Query 2

$$\sigma_{dept\_name = "Physics"} (\sigma_{salary > 90.000} (instructor))$$

- The two queries are not identical;
- they are however, equivalent –
  - they give the same result on any database.

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |
| 32343 | HIS-351   | 1      | Spring   | 2010 |
| 45565 | CS-101    | 1      | Spring   | 2010 |
| 45565 | CS-319    | 1      | Spring   | 2010 |
| 76766 | BIO-101   | 1      | Summer   | 2009 |
| 76766 | BIO-301   | 1      | Summer   | 2010 |
| 83821 | CS-190    | 1      | Spring   | 2009 |
| 83821 | CS-190    | 2      | Spring   | 2009 |
| 83821 | CS-319    | 2      | Spring   | 2010 |
| 98345 | EE-181    | 1      | Spring   | 2009 |

# Equivalent Queries

- There is more than one way to write a query in relational algebra.
- Another Example: Find information about courses taught by instructors in the Physics department
  - Query 1

$$\sigma_{dept\_name = "Physics"}(instructor \bowtie_{instructor.ID = teaches.ID} teaches)$$

- Query 2

$$(\sigma_{dept\_name = "Physics"}(instructor)) \bowtie_{instructor.ID = teaches.ID} teaches$$

- The two queries are not identical; they are, however, equivalent -- they give the same result on any database.

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 22222 | Einstein   | Physics    | 95000  |
| 12121 | Wu         | Finance    | 90000  |
| 32343 | El Said    | History    | 60000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 76766 | Crick      | Biology    | 72000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 58583 | Califieri  | History    | 62000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 76543 | Singh      | Finance    | 80000  |

| ID    | course_id | sec_id | semester | year |
|-------|-----------|--------|----------|------|
| 10101 | CS-101    | 1      | Fall     | 2009 |
| 10101 | CS-315    | 1      | Spring   | 2010 |
| 10101 | CS-347    | 1      | Fall     | 2009 |
| 12121 | FIN-201   | 1      | Spring   | 2010 |
| 15151 | MU-199    | 1      | Spring   | 2010 |
| 22222 | PHY-101   | 1      | Fall     | 2009 |
| 32343 | HIS-351   | 1      | Spring   | 2010 |
| 45565 | CS-101    | 1      | Spring   | 2010 |
| 45565 | CS-319    | 1      | Spring   | 2010 |
| 76766 | BIO-101   | 1      | Summer   | 2009 |
| 76766 | BIO-301   | 1      | Summer   | 2010 |
| 83821 | CS-190    | 1      | Spring   | 2009 |
| 83821 | CS-190    | 2      | Spring   | 2009 |
| 83821 | CS-319    | 2      | Spring   | 2010 |
| 98345 | EE-181    | 1      | Spring   | 2009 |

# Relational Algebra Exercises

*employee (person\_name, street, city)*

*works (person\_name, company\_name, salary)*

*company (company\_name, city)*

- Q1: Find the names of all employees who live in city “Miami”.

$$\Pi_{\text{person\_name}} (\sigma_{\text{city} = \text{“Miami”}} (\text{employee}))$$

- Q2: Find the names of all employees whose salary is greater than \$100,000.

$$\Pi_{\text{person\_name}} (\sigma_{\text{salary} > 100000} (\text{employee} \bowtie \text{works}))$$

- Q3: Find the names of all employees who live in “Miami” and whose salary is greater than \$100,000.

$$\Pi_{\text{person\_name}} (\sigma_{\text{city} = \text{“Miami”} \wedge \text{salary} > 100000} (\text{employee} \bowtie \text{works}))$$

# Relational Algebra Exercises

*branch(branch\_name, branch\_city, assets)*  
*customer (customer\_name, customer\_street, customer\_city)*  
*loan (loan\_number, branch\_name, amount)*  
*borrower (customer\_name, loan\_number)*  
*account (account\_number, branch\_name, balance)*  
*depositor (customer\_name, account\_number)*

- Q4: Find the names of all branches located in “Chicago”.

$$\Pi_{branch\_name} (\sigma_{branch\_city = \text{“Chicago”}} (branch))$$

- Q5: Find the names of all borrowers who have a loan in branch “Downtown”.

$$\Pi_{ID} (\sigma_{branch\_name = \text{“Downtown”}} (borrower \bowtie loan))$$

# Relational Algebra Exercises

PUBLISHER

| <u>publisher_id</u> | publisher_name | phone |
|---------------------|----------------|-------|
|---------------------|----------------|-------|

AUTHOR

| <u>author_id</u> | author_name | email |
|------------------|-------------|-------|
|------------------|-------------|-------|

BOOK

| <u>ISBN</u> | <u>publisher_id</u> | <u>author_id</u> | title |
|-------------|---------------------|------------------|-------|
|-------------|---------------------|------------------|-------|

**QUERY : Find titles of books that were published by the publisher id#20.**

**RA Query: Selection; Projection**

$$\pi \text{ title } (\sigma_{\text{publisher\_id} = 20} (\text{BOOK}))$$

# Relational Algebra Exercises

PUBLISHER

| <u>publisher_id</u> | publisher_name | phone |
|---------------------|----------------|-------|
|---------------------|----------------|-------|

AUTHOR

| <u>author_id</u> | author_name | email |
|------------------|-------------|-------|
|------------------|-------------|-------|

BOOK

| <u>ISBN</u> | <u>publisher_id</u> | <u>author_id</u> | title |
|-------------|---------------------|------------------|-------|
|-------------|---------------------|------------------|-------|

**QUERY : Find all book titles, along with their author name and publisher name.**

RA Query

$$\pi_{\text{title}, \text{author\_name}, \text{publisher\_name}} ( (\text{BOOK} \bowtie \text{AUTHOR} \bowtie \text{PUBLISHER}) )$$

# Summary of Relational Algebra Operators

| Symbol (Name)                   | Example of Use                                                                                                                                                                  |
|---------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\sigma$<br>(Selection)         | $\sigma \text{ salary} >= 85000 \text{ } (instructor)$<br>Return rows of the input relation that satisfy the predicate.                                                         |
| $\Pi$<br>(Projection)           | $\Pi ID, salary \text{ } (instructor)$<br>Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output.                             |
| $\times$<br>(Cartesian Product) | $instructor \times department$<br>Output pairs of rows from the two input relations (regardless of whether they have the same value on all attributes that have the same name). |
| $\cup$<br>(Union)               | $\Pi name \text{ } (instructor) \cup \Pi name \text{ } (student)$<br>Output the union of tuples from the <i>two</i> input relations.                                            |
| $-$<br>(Set Difference)         | $\Pi name \text{ } (instructor) -- \Pi name \text{ } (student)$<br>Output the set difference of tuples from the two input relations.                                            |
| $\bowtie$<br>(Natural Join)     | $instructor \bowtie department$<br>Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.                        |