



Istanbul Technical University  
Department of Computer Engineering

## BLG 322E - COMPUTER ARCHITECTURE

### Assignment 4

**Due Date: Wednesday, May 23, 2025, 23:59.**

- Please **write your full name** (first name and last name) **and Student ID** at the top of your solution.
- Please show ALL your work. Answers with no supporting explanations or work will not receive any partial credit. Your homework is not just a final report of your results; we want to see your steps. Upload all the papers you worked on to get to the solution.
- Please follow all formatting instructions exactly (tables, address notation, naming conventions, etc.). Use the assignment file (.docx) directly by filling in blank (red) fields or copy them faithfully in your own report.
- **Submissions:** Submit your solution as a PDF file to Ninova before the deadline.
- **No late submissions** will be accepted. Do not send your solutions by e-mail. We will only accept files that have been uploaded to the official Ninova e-learning system before the deadline. Do not risk leaving your submission to the last few minutes.
- **Consequences of plagiarism/cheating:** Assignments have to be done individually. Any cheating will be subject to disciplinary action.
- **AI usage:** Language and reasoning models (ChatGPT, o1, o3, Claude, etc.) are tested on the questions. They hallucinate on some aspects of questions and give the wrong answers. You can use them to test alternative ideas, but you should always check the results.

If you have any questions, please e-mail **Mustafa Abdullah Hakkoz** ([hakkoz22@itu.edu.tr](mailto:hakkoz22@itu.edu.tr)) or use the message board in Ninova.

Student Name: Mustafa Can Çalışkan  
Student No: 150200097

## Assignment 4: Cache-Memory Analysis and Access Time

A computer system has the following characteristics:

| Parameter         | Value  |
|-------------------|--|
| Main memory       | <b>512 KByte</b> (19-bit physical addresses)                                 |
| L1 data-cache     | <b>512 bytes, 4-way set-associative</b>                                      |
| Block (line) size | <b>16 bytes</b>  |
| Replacement       | <b>LRU</b> with a 2-bit ageing counter<br>(00 = oldest ... 11 = most-recent) |
| Read policy       | Read-through   |
| Write policy      | Write-through + Write-allocate   |
| Transfers         | Whole 16-byte blocks move between memory and cache                           |

The cache is **initially empty**. The CPU executes the code below.

**Program in c (outer loop runs exactly two times)**

```
for (int i = 1; i <= 2; ++i)
{
    /* 1. A - forward scan (16 elements)*/
    for (int j = 0; j < 16; ++j)    // A[0]...A[15]
        t += A[j];

    /* 2. B - backward scan (16 elements)*/
    for (int j = 15; j >= 0; --j)    // B[15]...B[0]
        t += B[j];

    /* 3. C - forward scan (24 elements) */
    for (int j = 0; j < 24; ++j)    // C[0]...C[23]
        t += C[j];

    /* 4. D - ping-pong scan, starts->both sides, ends->centre (32 elements) */
    for (int j = 0; j < 16; ++j) {
        t += D[j];
        t += D[31-j];
    }

    /* 5. E - stride-2 scan (48 elements)*/
    for (int j = 0; j < 48; j += 2)    // E[0],E[2]...E[46]
        t += E[j];
}
```

| Array | Base address | Size |
|-------|--------------|------|
| A     | 0x0000 A0    | 16 B |
| B     | 0x0040 20    | 16 B |
| C     | 0x0150 10    | 24 B |
| D     | 0x0200 90    | 32 B |
| E     | 0x0280 80    | 48 B |

All elements are **one byte**; the loops **read only**.

Use timing symbols:

- $t_c$  – cache-hit time
  - $t_m$  – main-memory access time
  - $t_b$  – block-transfer time
- $$t_b = 2t_m = 10t_c,$$

$$t_b > 0$$

## Tasks

- a) **(5 pts)** Into which fields does the cache controller divide the physical address? Specify field names and bit widths (Tag, Set Index, Block Offset/Word bits).

**Table-1: Address-field breakdown**

| Field                       | Bits | Reason   |
|-----------------------------|------|--|
| Block Offset<br>(Word bits) | 4    | The block line size is 16. 4 bits are enough to represent offset.  |
| Set Index                   | 3    | $512/16 = 32$ frame. Each frame consists of 4 blocks (4-way). $32/4 = 8$ , meaning that 3 bits are enough to represent set offset. |
| Tag                         | 12   | Physical address size – block offset – set index = remaining = 12.   |

- b) **(5 pts)** Describe the **tag-store memory**:

- number of sets (rows)
- number of ways (frames) per set
- tag width per way
- total tag-memory size in **bytes** (include status bits: 1 valid bit, 1 dirty bit, and 2 LRU-ageing bits per frame).

**Table-2: Tag-memory size calculation**

| Item                  | Calculation  |
|-----------------------|--|
| Sets                  | $512/16 = 32$ frame. Each frame consists of 4 blocks (4-way). $32/4 = 8$ . There are 8 sets. |
| Ways (frames) per set | 4-ways. 4 frame. It is given.  |
| Tag bits per frame    | 12 bits. Found in part a).   |
| Status bits per frame | 1 valid bit + 1 dirty bit + 2 aging bits = 4 bits.   |
| Bits per frame        | 12 (tag bits) + 4 (status bits) = 16.  |
| Bits per row          | $16 * 4$ (4-way) = 64. I assume that row indicates set.                                      |
| Bytes per row         | $64 / 8 = 8$ bytes.  |
| Total tag memory      | 8 sets * 8 bytes = 64 bytes.   |

**c) (30 pts)** For the **first outer-loop iteration only** ( $i = 1$ ), show **which memory blocks map to which set and way** after the entire iteration completes. Indicate **misses, hits, and replacements** by filling following table:

**Table-3: Distinct blocks** (example row for A already filled)

| Array    | Block label | Hex base  | 19-bit Tag   Set Idx   Word Number | Set # | 1st byte of the block (Hit or Miss) |
|----------|-------------|-----------|------------------------------------|-------|-------------------------------------|
| <b>A</b> | $A_0$       | 0x0000 A0 | 000 0000 0000 1   010   0000       | 2     | M                                   |
| <b>B</b> | $B_0$       | 0x0040 20 | 000 0100 0000 0   010   0000       | 2     | H                                   |
| <b>C</b> | $C_0$       | 0x015010  | 001 0101 0000 0   001   0000       | 1     | M                                   |

|          |       |          |                              |   |   |
|----------|-------|----------|------------------------------|---|---|
| <b>C</b> | $C_1$ | 0x015020 | 001 0101 0000 0   010   0000 | 2 | M |
| <b>D</b> | $D_0$ | 0x020090 | 010 0000 0000 1   001   0000 | 1 | M |
| <b>D</b> | $D_1$ | 0x0200A0 | 010 0000 0000 1   010   0000 | 2 | H |
| <b>E</b> | $E_0$ | 0x028080 | 010 1000 0000 1   000   0000 | 0 | M |
| <b>E</b> | $E_1$ | 0x028090 | 010 1000 0000 1   001   0000 | 1 | M |
| <b>E</b> | $E_2$ | 0x0280A0 | 010 1000 0000 1   010   0000 | 2 | M |

**NOTE:** A single row in the table represents one 16-byte cache block. If an array spans several blocks, just insert additional rows to list those blocks. The example of array A is 16 bytes long, so it fits entirely in one block ( $A_0$ ) and needs only that single row. If A were larger than 16 bytes you would add rows  $A_1$ ,  $A_2$ , and so on for each extra block.

After filling table 3, illustrate the evolving LRU order for any set that receives more than four distinct blocks. Provide the **contents of the tag-store rows** you touched (tag+valid+dirty+ageing counter) at the end of this first iteration by filling following table:

**Table-4: Cache snapshot after pass 1 ( $i=1$ )** (example row shows the format for set 0)

| Set      | Way0<br>Tag(3) in hex**<br>  V(1)D(1)Age(2) bits | Way1<br>Tag(3) in hex<br>  V(1)D(1)Age(2) bits | Way2<br>Tag(3) in hex<br>  V(1)D(1)Age(2) bits | Way3<br>Tag(3) in hex<br>  V(1)D(1)Age(2) bits | Removed<br>blocks* |
|----------|--|--|--|--|--------------------|
| <b>0</b> | 0x501   1011 ( $E_0$ )                           | —  | —  | —  | —                  |
| <b>1</b> | 0x2A0   1001 ( $C_0$ )                           | 0x401   1010 ( $D_0$ )                         | 0x501   1011 ( $E_1$ )                         | -  | -                  |
| <b>2</b> | 0x080   1000 ( $B_0$ )                           | 0x2A0   1001 ( $C_1$ )                         | 0x401   1010 ( $D_1$ )                         | 0x501   1011 ( $E_2$ )                         | A0                 |

\* Please write tag bits in hexadecimal and status bits (valid(1), dirty(1), ageing(2)) in binary. Use 11 to mark the most-recently-used (MRU) frame and 00 for the least-recently-used (LRU) frame.

\*\* Blocks removed from the row of sets to make space for new ones. For the case of LRU, it is in the least used frame. Indicate them here.

d) (20 pts) Now let the outer loop finish all two passes.

- i. How many **total replacements** occur during the entire run? Fill the table below and show your work.
- ii. Compute the **overall hit ratio**. Calculate hits and misses on byte level and fill the table. Remember when you check for miss, you only need to look at first byte of the block. On the other hand, for counting hits, you need to look at all bytes of the block.

**Table-5: Hit/Miss calculations**

| Metric            | After pass 1 (i==1) | After pass 2 (i==2) | Total |
|-------------------|---------------------|---------------------|-------|
| Byte hits         | 103                 | 107                 | 210   |
| Byte misses       | 9                   | 5                   | 14    |
| Replacements      | 1 block (16 bytes)  | 5 block (80 bytes)  | 6     |
| Hit ratio (bytes) | 0.92                | 0.96                | 0.94  |

e) (10 pts) Propose a **new starting address for array E** (expressed in hexadecimal) that **reduces conflict misses** without changing any of the other addresses. Explain concisely **why** your choice helps.

**Answer for part e:**

New start address: 0x028030. Reason: At this address, the set number starts with 3, and when the entire array is fetched, it is distributed across sets 3, 4, and 5. This way, since no more than 4 blocks will be mapped to Set 2, the extreme conflict is resolved.

f) (10 pts) Using your miss count for pass 1 in **part c**, compute the **stall time**. Express the result in  $t_c$ .

$$T_1 = misses(t_m + t_b) + hits(t_c)$$

Answer for part f:

$$T1 = 9(5tc+10tc)+103tc=9(15tc)+103tc=238tc$$

**g) (10 pts)** Repeat the calculation after pass 2 (total time passed) and derive the **average time for memory access**.

$$Av. time = \frac{total\ time}{number\ of\ total\ byte\ accesses}$$

Answer for part g:

$$Total\ time = 14(5tc+10tc)+210tc=420tc$$

$$Number\ of\ total\ byte\ access = 224$$

$$Av.\ Time = 1.88tc$$

**h) (10 pts)** Explain briefly why a FIFO policy would or would not change your numerical average time for this specific program.

Answer for part h:

If FIFO had been used, the result would not change. This is because each array is accessed only once and in order within the loop (in the order A B C D E), and since the same array is not accessed again within the same loop (the aging counters of the arrays decrease (getting older) according to their access order), the block that entered the cache first (the oldest one in LRU) is always evicted when replacement is occurred. This causes LRU to behave like FIFO in this code.